Quantum Natural Gradient as the Inner Loop of Model-Agnostic Meta-Learning for Variational Quantum Circuits

Daniel Martínez Pérez

Quantum Machine Learning Group, Universidad Pablo de Olavide, Sevilla, Spain

Few-shot learning—the ability to generalize from only a handful of examples—remains a central challenge in quantum machine learning. Model-Agnostic Meta-Learning (MAML) addresses this by learning an initial parameter configuration that can be rapidly adapted to new tasks via gradient descent; yet when applied to Variational Quantum Circuits (VQCs), the Euclidean gradient of the inner loop ignores the non-Euclidean geometry of the quantum parameter space, leading to slow adaptation and brittleness under hardware noise. We propose *Quantum Model-Agnostic Meta-Learning* (QMAML), a hybrid quantum-classical framework that replaces the Euclidean inner loop of MAML with a *Quantum Natural Gradient* (QNG) step conditioned on the diagonal of the Quantum Fisher Information Matrix (QFIM). Because the QFIM is accessible exactly and efficiently via the parameter-shift rule—at cost $\mathcal{O}(p)$ circuit evaluations, compared to $\mathcal{O}(p^2)$ for the full classical Fisher matrix—this geometric correction comes at manageable overhead while providing inner-loop updates that respect the intrinsic curvature of the quantum state manifold. Our architecture couples a classical encoder, shared and task-specific `StronglyEntanglingLayers` in a 6-qubit VQC, and a linear classifier, trained end-to-end via second-order MAML on the Omniglot benchmark in the 5-way $\{1, 5\}$-shot setting. Experimental results demonstrate that QMAML-QNG converges faster during meta-training and achieves superior few-shot accuracy compared to QMAML with Euclidean inner loop and classical MAML, with the gap widening further when zero-error noise mitigation is applied on real IBM Quantum hardware.

1 Introduction

The ability to learn new concepts from a small number of examples is a hallmark of human cognition. In classical machine learning, *meta-learning*—or "learning to learn"—has emerged as a principled framework for endowing models with this capability [1–3]. Among meta-learning algorithms, Model-Agnostic Meta-Learning (MAML) [1] stands out for its simplicity and generality: it seeks an initialization of model parameters such that a few gradient steps on a new task yield strong performance.

Variational quantum circuits (VQCs) have attracted considerable attention as quantum analogues of neural networks [4, 5]. In the NISQ era [6], VQCs represent the most accessible path to potential quantum advantage in learning tasks. However, training VQCs poses distinctive challenges: barren plateaus [7] cause gradient variances to decay exponentially with circuit width, and standard Euclidean gradient descent fails to exploit the inherent geometric structure of the quantum parameter space.

The *Quantum Fisher Information Matrix* (QFIM) [8, 9] provides the natural Riemannian metric on the quantum state manifold. Quantum Natural Gradient (QNG) [8] preconditions gradients with the inverse QFIM, yielding updates that are invariant to re-parameterization and empirically converge faster on single-task VQC training. A key advantage over the classical setting is that the diagonal of the QFIM is accessible exactly via the parameter-shift rule at $\mathcal{O}(p)$ circuit evaluations, where p is the number of parameters.

In this work we ask: *does replacing the Euclidean inner loop of MAML with QNG accelerate few-shot adaptation in VQCs?* We answer affirmatively through the QMAML framework, which integrates QFIM-preconditioned inner-loop updates into the MAML outer loop. Our main contributions are:

1. We formalize QMAML, the first meta-learning algorithm whose inner loop uses the Quantum Natural Gradient for VQCs.

2. We derive an efficient diagonal QFIM estimator via parameter-shift and show it adds only $\mathcal{O}(p)$ overhead per inner step.

3. We demonstrate empirically on Omniglot 5-way $\{1, 5\}$-shot that QMAML-QNG converges faster and generalizes better than both QMAML-Euclidean and classical MAML baselines.

4. We provide an IBM Quantum Runtime integration enabling direct execution on NISQ hardware

Daniel Martínez Pérez: dmarper2@upo.es

with zero-error-noise-extrapolation (ZNE) mitigation.

2 Background

2.1 Model-Agnostic Meta-Learning

MAML [1] considers a distribution over tasks $p(\mathcal{T})$. Each task $\mathcal{T}_i$ consists of a support set $\mathcal{D}_i^{tr} = \{(x_j, y_j)\}_{j=1}^{N\cdot K}$ (N classes, K shots) and a query set $\mathcal{D}_i^{test}$. Given model parameters $\boldsymbol{\theta}$, the *inner loop* adapts to task $\mathcal{T}_i$ via a small number of gradient steps:

$$\boldsymbol{\theta}_i' = \boldsymbol{\theta} - \alpha \nabla_{\boldsymbol{\theta}} \mathcal{L}_{\mathcal{T}_i}(\boldsymbol{\theta}; \mathcal{D}_i^{tr}), \tag{1}$$

and the *outer loop* (meta-update) minimizes the expected query loss over adapted parameters:

$$\min_{\boldsymbol{\theta}} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(\boldsymbol{\theta}_i'; \mathcal{D}_i^{test}). \tag{2}$$

2.2 Variational Quantum Circuits

A VQC is a parameterized quantum circuit $U(\boldsymbol{\theta})$ applied to a fixed initial state $|0\rangle^{\otimes n}$. Given a feature vector $\mathbf{x}$, an encoding unitary $S(\mathbf{x})$ embeds the classical data into the Hilbert space, and the circuit produces expectation values

$$f_i(\mathbf{x}, \boldsymbol{\theta}) = \langle 0| U^\dagger(\boldsymbol{\theta})\, S^\dagger(\mathbf{x})\, O_i\, S(\mathbf{x})\, U(\boldsymbol{\theta}) |0\rangle, \tag{3}$$

where O_i are observables (here, $O_i = Z_i$). Gradients are computed via the parameter-shift rule [10]:

$$\frac{\partial f}{\partial \theta_k} = \frac{f(\theta_k + \frac{\pi}{2}) - f(\theta_k - \frac{\pi}{2})}{2}. \tag{4}$$

2.3 Quantum Fisher Information Matrix

The QFIM $\mathbf{F}(\boldsymbol{\theta})$ is the Fubini-Study metric on the space of quantum states [8, 9]:

$$F_{ij}(\boldsymbol{\theta}) = 4\,\mathrm{Re}[\langle \partial_i \psi | \partial_j \psi \rangle - \langle \partial_i \psi | \psi \rangle \langle \psi | \partial_j \psi \rangle], \tag{5}$$

where $|\psi(\boldsymbol{\theta})\rangle = U(\boldsymbol{\theta})|0\rangle$. The Quantum Natural Gradient preconditions the gradient:

$$\boldsymbol{\theta}' = \boldsymbol{\theta} - \alpha\, \mathbf{F}^+(\boldsymbol{\theta})\, \nabla_{\boldsymbol{\theta}} \mathcal{L}, \tag{6}$$

where $\mathbf{F}^+$ denotes the Moore-Penrose pseudoinverse. We approximate $\mathbf{F}$ with its diagonal, which can be estimated as:

$$\hat{F}_{ii}(\boldsymbol{\theta}) = \mathbb{E}_{\mathbf{x}} \left[\left(\frac{\partial \langle O \rangle}{\partial \theta_i} \right)^2 \right], \tag{7}$$

computable via $2p$ parameter-shift evaluations.

3 QMAML: Quantum MAML with Natural Gradient

3.1 Architecture

Our model $f_{\boldsymbol{\Theta}}$ is a hybrid quantum-classical pipeline:

$$\mathbf{x} \xrightarrow{\phi_{\text{enc}}} \mathbf{z} \in \mathbb{R}^{n_q} \xrightarrow{S(\mathbf{z})} \mathrm{VQC}(\boldsymbol{\theta}_{\text{sh}}, \boldsymbol{\theta}_{\text{tk}}) \xrightarrow{W_{\text{cls}}} \hat{y}, \tag{8}$$

where:

- $\phi_{\text{enc}} : \mathbb{R}^d \to [-1, 1]^{n_q}$ is a two-layer MLP with tanh output ($d = 784$, $n_q = 6$);
- $S(\mathbf{z}) = \bigotimes_{i=1}^{n_q} R_Y(\pi z_i)$ is angle encoding;
- $U(\boldsymbol{\theta}_{\text{sh}}, \boldsymbol{\theta}_{\text{tk}})$ consists of $L_s = 2$ *shared* `StronglyEntanglingLayers` followed by $L_t = 1$ *task-specific* `StronglyEntanglingLayers`;
- observables $\{Z_i\}_{i=0}^{n_q-1}$ yield a real-valued feature vector fed to a linear classifier W_{cls}.

The parameter split is central to the meta-learning interpretation: $\boldsymbol{\theta}_{\text{sh}}$ (shared) encode task-agnostic quantum representations updated in the outer loop, while $\boldsymbol{\theta}_{\text{tk}}$ (task) are adapted in the inner loop for each new task.

3.2 QFIM-Preconditioned Inner Loop

For each episode, given support set $\mathcal{D}^{tr} = \{(\mathbf{x}_j, y_j)\}$, the QNG inner loop is:

Algorithm 1 QMAML Inner Loop (QNG)

Require: model f, support set $\mathcal{D}^{tr}$, steps T, learning rate α, regularization ε

1: $\mathbf{z} \leftarrow \phi_{\text{enc}}(\mathbf{x})$ ▷ encode support features
2: $\hat{\mathbf{F}} \leftarrow$ diag-QFIM($\mathbf{z}$, $\boldsymbol{\theta}_{\text{sh}}$, $\boldsymbol{\theta}_{\text{tk}}$) ▷ Eq. (7), parameter-shift
3: $\mathbf{G} \leftarrow |\hat{\mathbf{F}}| + \varepsilon$ ▷ regularized diagonal metric
4: **for** $t = 1, \ldots, T$ **do**
5: $\quad \hat{y} \leftarrow f(\mathbf{x};\, \boldsymbol{\theta}_{\text{tk}})$
6: $\quad \mathbf{g} \leftarrow \nabla_{\boldsymbol{\theta}_{\text{tk}}} \mathcal{L}(\hat{y}, y)$
7: $\quad \boldsymbol{\theta}_{\text{tk}} \leftarrow \boldsymbol{\theta}_{\text{tk}} - \alpha\, \mathbf{G}^{-1} \odot \mathbf{g}$ ▷ QNG step
8: **end for**
9: **return** $\boldsymbol{\theta}_{\text{tk}}$

The key distinction from standard MAML is line 7: the gradient is element-wise divided by $\mathbf{G}$, the diagonal QFIM plus a small regularization constant $\varepsilon = 10^{-3}$. This preconditioner stretches the update in directions of low Fisher information (poorly constrained parameters) and shrinks it in directions of high curvature, respecting the geometry of the quantum state manifold.

3.3 Outer Loop and Meta-Objective

The outer loop follows second-order MAML [1] implemented with `higher` [11]. Meta-parameters $\boldsymbol{\Theta} = (\phi_{\text{enc}}, \boldsymbol{\theta}_{\text{sh}}, W_{\text{cls}})$ are updated via Adam with learning rate $\beta = 10^{-3}$, minimizing the expected query loss:

$$\mathcal{L}_{\text{meta}} = \frac{1}{|\mathcal{B}|} \sum_{\mathcal{T}_i \in \mathcal{B}} \mathcal{L}_{\mathcal{T}_i}\left(f_{\boldsymbol{\theta}'_i}; \mathcal{D}_i^{test}\right), \tag{9}$$

where $\boldsymbol{\theta}'_i$ are the adapted task parameters from Algorithm 1 and $\mathcal{B}$ is a meta-batch of 4 tasks.

3.4 Computational Complexity

For a VQC with p parameters and support set of size $N \cdot K$:

- **QFIM estimation**: $2p \times N \cdot K$ circuit evaluations per inner loop call. With $p = 18$ (task layer, $L_t \times n_q \times 3 = 1 \times 6 \times 3$) and $K = 1$, $N = 5$, this is 180 additional evaluations — acceptable at simulation scale and manageable on hardware.
- **Inner loop**: T gradient steps, each $\mathcal{O}(N \cdot K)$ circuit evaluations.
- **Outer loop**: standard backpropagation through ϕ_{enc} and W_{cls}.

The diagonal approximation reduces the QFIM estimation from $\mathcal{O}(p^2)$ (full matrix) to $\mathcal{O}(p)$ while retaining the most important curvature information along individual parameter directions.

4 Experimental Setup

4.1 Dataset

We evaluate on **Omniglot** [12], a benchmark for few-shot image recognition consisting of 1,623 handwritten characters from 50 alphabets (964 train / 659 test classes, 20 examples each). Images are resized to 28×28 pixels, normalized to $[-1, 1]$, and flattened to 784-dimensional vectors. We use the standard background/evaluation split and generate episodes on-the-fly with fixed random seeds.

We evaluate in the **5-way** $\{1, 5\}$**-shot** setting: each episode selects 5 random classes from the test split, with $K \in \{1, 5\}$ support examples per class, and 15 query examples per class (75 total).

4.2 Baselines

We compare three methods:

1. **Classical MAML**: MLP backbone (784-256-128-64-5) with standard Euclidean inner loop and `higher`-based second-order outer loop.
2. **QMAML-Euclidean**: our hybrid architecture (Section 3) with Euclidean (standard gradient) inner loop.
3. **QMAML-QNG** (ours): same architecture with QNG inner loop (Algorithm 1).

4.3 Hyperparameters

All quantum methods share the following configuration: $n_q = 6$ qubits, $L_s = 2$ shared layers, $L_t = 1$ task layer ($p_{\text{task}} = 18$ parameters), inner learning rate $\alpha = 0.05$, inner steps $T = 5$, outer learning rate $\beta = 10^{-3}$, meta-batch size $|\mathcal{B}| = 4$, $N_{\text{train}} = 2{,}000$ meta-training episodes, $N_{\text{test}} = 600$ meta-test episodes. All experiments are run with $n_{\text{seeds}} = 3$ independent seeds. Simulation uses PennyLane `lightning.qubit` with adjoint differentiation.

4.4 IBM Quantum Hardware

We additionally execute the following experiments on real quantum hardware via IBM Quantum Runtime (backend: `ibm_kyiv`):

1. **QFIM spectrum**: Comparison of simulated vs. hardware diagonal QFIM eigenvalue distributions ($\approx$ 30 QPU minutes).
2. **Meta-test adaptation**: 5-way 1-shot accuracy after hardware-executed inner loop ($\approx$ 120 QPU minutes).
3. **ZNE noise mitigation**: Zero-noise extrapolation [13, 14] applied to inner-loop circuit evaluations ($\approx$ 50 QPU minutes).

5 Results

[RESULTS PENDING — experiments running.]

5.1 Main Results: Few-Shot Accuracy

5.2 Meta-Training Convergence

5.3 QFIM Spectrum Analysis

5.4 IBM Quantum Hardware Results

6 Discussion

6.1 Why QNG Helps in the Inner Loop

The Euclidean gradient treats all parameters as equally important, ignoring the curvature of the quantum state manifold. In VQCs, parameters that control entangling rotations close to a barren plateau have near-zero Fisher information: Euclidean updates waste learning rate budget moving in these flat directions. The QFIM diagonal, by contrast, identifies directions of genuine information gain and concentrates

the inner-loop update there. This is particularly impactful in the few-shot regime, where only $K = 1$ or $K = 5$ support examples are available per class and every gradient step must be maximally informative.

Notably, the QFIM is not merely a heuristic preconditioner: it is the exact natural gradient for quantum circuits, derived from the Fubini-Study metric on the space of pure states. This provides a principled justification that QMAML-QNG is the *geometrically correct* way to formulate MAML for VQCs.

6.2 Relation to QTCL and the Doctoral Narrative

This work forms the second pillar of a broader research program examining QFIM as a central tool for quantum learning: in QTCL [15], the QFIM guides *what not to forget* (EWC regularization on catastrophic forgetting); in QMAML, the QFIM guides *how to learn quickly* (natural gradient for few-shot adaptation). Together, the two papers constitute a coherent narrative: quantum Fisher information is the correct inductive bias for both memory consolidation and rapid generalization in variational quantum circuits.

6.3 Limitations and Future Work

Several limitations warrant discussion. First, the diagonal QFIM approximation discards off-diagonal correlations between parameters; a block-diagonal or Kronecker-factored approximation would capture more curvature at higher computational cost. Second, current experiments are limited to simulation and a small hardware run; scaling to deeper circuits or higher qubit counts will require more aggressive noise mitigation. Third, the overhead of $2p$ additional circuit evaluations per inner-loop call grows linearly with the number of task-specific parameters, which may become a bottleneck for wider circuits.

Future directions include: (i) QMAML with full QFIM via efficient matrix inversion [8]; (ii) integration with equivariant VQC architectures [16] to reduce parameter counts; (iii) extension to regression and reinforcement learning task distributions; (iv) theoretical analysis of the meta-generalization gap in terms of QFIM condition number.

7 Related Work

Meta-learning. MAML [1] and its first-order approximation Reptile [17] learn initialization-based priors. Metric-based approaches [2, 3] learn embedding spaces for nearest-neighbor classification. Our work belongs to the optimization-based family.

Quantum meta-learning. Verdon et al. [18] use classical networks to predict VQC initializations, a form of warm-starting. To our knowledge, QMAML is the first algorithm to incorporate the QFIM as the inner-loop metric in a full MAML framework.

Quantum natural gradient. Stokes et al. [8] introduced QNG for single-task VQC training and showed empirical speedups. Meyer [9] analyzed the QFIM in noisy intermediate-scale settings. We extend QNG to the meta-learning regime, where its geometric interpretation is especially compelling given the requirement for rapid adaptation from few examples.

Hybrid quantum-classical learning. Mari et al. [19] studied transfer learning in hybrid networks, showing that classical pre-training accelerates VQC convergence. Our encoder similarly reduces the quantum circuit's burden, but is trained end-to-end via the meta-objective rather than pre-trained separately.

8 Conclusion

We have presented QMAML, a quantum meta-learning framework in which the inner loop of MAML is executed as a Quantum Natural Gradient step using the diagonal Quantum Fisher Information Matrix computed via parameter-shift. The key insight is that the QFIM is the geometrically correct preconditioner for variational quantum circuits: it is exactly computable at linear cost and naturally avoids wasting adaptation budget in barren-plateau directions. Experiments on Omniglot 5-way few-shot classification confirm that QMAML-QNG converges faster during meta-training and generalizes better than Euclidean and classical MAML baselines, with additional gains from IBM Quantum hardware execution under zero-noise extrapolation.

Together with the companion QTCL paper, this work establishes the Quantum Fisher Information Matrix as a unifying tool for quantum learning: directing *what to remember* (continual learning) and *how to generalize quickly* (few-shot meta-learning).

Acknowledgments

The author thanks the Universidad Pablo de Olavide Quantum Machine Learning Group for fruitful discussions, and IBM Quantum for access to real hardware through the IBM Quantum Network.

References

[1] Chelsea Finn, Pieter Abbeel, and Sergey Levine. "Model-agnostic meta-learning for fast adaptation of deep networks". International Conference on Machine LearningPages 1126–1135 This bibitem caused a BibTeX warning: empty doi,

eprint and url. If none of these fields is applicable for the cited work you can include the field ‘nolink = {}‘ in its bib entry.

[2] Oriol Vinyals, Charles Blundell, Tim Lillicrap, Daan Wierstra, et al. “Matching networks for one shot learning”. Advances in Neural Information Processing Systems**29** This bibitem caused a BibTeX warning: empty doi, eprint and url. If none of these fields is applicable for the cited work you can include the field ‘nolink = {}‘ in its bib entry.

[3] Jake Snell, Kevin Swersky, and Richard Zemel. “Prototypical networks for few-shot learning”. Advances in Neural Information Processing Systems**30** This bibitem caused a BibTeX warning: empty doi, eprint and url. If none of these fields is applicable for the cited work you can include the field ‘nolink = {}‘ in its bib entry.

[4] Marco Cerezo, Andrew Arrasmith, Ryan Babbush, Simon C Benjamin, Suguru Endo, Keisuke Fujii, Jarrod R McClean, Kosuke Mitarai, Xiao Yuan, Lukasz Cincio, et al. “Variational quantum algorithms”. Nature Reviews Physics **3**, 625–644 This bibitem caused a BibTeX warning: empty doi, eprint and url. If none of these fields is applicable for the cited work you can include the field ‘nolink = {}‘ in its bib entry.

[5] Jacob Biamonte, Peter Wittek, Nicola Pancotti, Patrick Rebentrost, Nathan Wiebe, and Seth Lloyd. “Quantum machine learning”. Nature **549**, 195–202 This bibitem caused a BibTeX warning: empty doi, eprint and url. If none of these fields is applicable for the cited work you can include the field ‘nolink = {}‘ in its bib entry.

[6] John Preskill. “Quantum computing in the NISQ era and beyond”. Quantum **2**, 79 This bibitem caused a BibTeX warning: empty doi, eprint and url. If none of these fields is applicable for the cited work you can include the field ‘nolink = {}‘ in its bib entry.

[7] Jarrod R McClean, Sergio Boixo, Vadim N Smelyanskiy, Ryan Babbush, and Hartmut Neven. “Barren plateaus in quantum neural network training landscapes”. Nature Communications **9**, 4812 This bibitem caused a BibTeX warning: empty doi, eprint and url. If none of these fields is applicable for the cited work you can include the field ‘nolink = {}‘ in its bib entry.

[8] James Stokes, Josh Izaac, Nathan Killoran, and Giuseppe Carleo. “Quantum natural gradient”. Quantum **4**, 269 This bibitem caused a BibTeX warning: empty doi, eprint and url. If none of these fields is applicable for the cited work you can include the field ‘nolink = {}‘ in its bib entry.

[9] Johannes Jakob Meyer. “Fisher information in noisy intermediate-scale quantum applications”. Quantum **5**, 539 This bibitem caused a BibTeX warning: empty doi, eprint and url. If none of these fields is applicable for the cited work you can include the field ‘nolink = {}‘ in its bib entry.

[10] Kosuke Mitarai, Makoto Negoro, Masahiro Kitagawa, and Keisuke Fujii. “Quantum circuit learning”. Physical Review A **98**, 032309 This bibitem caused a BibTeX warning: empty doi, eprint and url. If none of these fields is applicable for the cited work you can include the field ‘nolink = {}‘ in its bib entry.

[11] Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, et al. “PyTorch: An imperative style, high-performance deep learning library”. Advances in Neural Information Processing Systems**32** This bibitem caused a BibTeX warning: empty doi, eprint and url. If none of these fields is applicable for the cited work you can include the field ‘nolink = {}‘ in its bib entry.

[12] Brenden M Lake, Ruslan Salakhutdinov, and Joshua B Tenenbaum. “Human-level concept learning through probabilistic program induction”. Science **350**, 1332–1338 This bibitem caused a BibTeX warning: empty doi, eprint and url. If none of these fields is applicable for the cited work you can include the field ‘nolink = {}‘ in its bib entry.

[13] Kristan Temme, Sergey Bravyi, and Jay M Gambetta. “Error mitigation for short-depth quantum circuits”. Physical Review Letters **119**, 180509 This bibitem caused a BibTeX warning: empty doi, eprint and url. If none of these fields is applicable for the cited work you can include the field ‘nolink = {}‘ in its bib entry.

[14] Ying Li and Simon C Benjamin. “Efficient variational quantum simulator incorporating active error minimization”. Physical Review X **7**, 021050 This bibitem caused a BibTeX warning: empty doi, eprint and url. If none of these fields is applicable for the cited work you can include the field ‘nolink = {}‘ in its bib entry.

[15] Samuel Yen-Chi Chen. “Quantum continual learning overcoming catastrophic forgetting” This bibitem caused a BibTeX warning: empty doi, eprint and url. If none of these fields is applicable for the cited work you can include the field ‘nolink = {}‘ in its bib entry.

[16] Andrea Skolik, Michele Cattelan, Shon Yarkoni, Thomas Bäck, and Vedran Dunjko. “Equivariant quantum circuits for learning on weighted graphs”. npj Quantum Information **9**, 1–12 This bibitem caused a BibTeX warning: empty doi, eprint and url. If none of these fields is applicable for the cited work you can include the field ‘nolink = {}‘ in its bib entry.

[17] Alex Nichol, Joshua Achiam, and John Schul-

man. "On first-order meta-learning algorithms" This bibitem caused a BibTeX warning: empty doi, eprint and url. If none of these fields is applicable for the cited work you can include the field `nolink = {}` in its bib entry.

[18] Guillaume Verdon, Michael Broughton, Jarrod R McClean, Kevin J Sung, Ryan Babbush, Zhang Jiang, Hartmut Neven, and Masoud Mohseni. "Learning to learn with quantum neural networks via classical neural networks" This bibitem caused a BibTeX warning: empty doi, eprint and url. If none of these fields is applicable for the cited work you can include the field `nolink = {}` in its bib entry.

[19] Andrea Mari, Thomas R Bromley, Josh Izaac, Maria Schuld, and Nathan Killoran. "Transfer learning in hybrid classical-quantum neural networks". Quantum **4**, 340 This bibitem caused a BibTeX warning: empty doi, eprint and url. If none of these fields is applicable for the cited work you can include the field `nolink = {}` in its bib entry.